

Final Project: Regression and Classification with Multi Architecture (Deep) Neural Networks

EMMANUEL OPOKU(22630361)

REPRODUCIBILITY

LIBRARIES

Introduction

The current research analyses models based on the following types of neural networks: deep neural networks (DNNs), convolutional neural networks (CNNs), long short-term memory networks (LSTMs), and hybrid networks. For regression analysis, the study utilises the Jena Climate dataset, while for the image classification task, the MNIST database will be used. The overall goal of this investigation is to analyse the effect of certain architectural features of neural networks on their ability to learn.

Part 0: Data Preparation and Exploratory Analysis

```
'data.frame':  420551 obs. of  15 variables:
 $ Date.Time      : chr  "01.01.2009 00:10:00" "01.01.2009 00:20:00" "01.01.2009 00:30:00" "
 $ p..mbar.       : num  997 997 997 997 997 ...
 $ T..degC.       : num  -8.02 -8.41 -8.51 -8.31 -8.27 -8.05 -7.62 -7.62 -
7.91 -8.43 ...
 $ Tpot..K.       : num  265 265 265 265 265 ...
 $ Tdew..degC.    : num  -8.9 -9.28 -9.31 -9.07 -9.04 -8.78 -8.3 -8.36 -
8.73 -9.34 ...
 $ rh....        : num  93.3 93.4 93.9 94.2 94.1 94.4 94.8 94.4 93.8 93.1 ...
 $ VPmax..mbar.   : num  3.33 3.23 3.21 3.26 3.27 3.33 3.44 3.44 3.36 3.23 ...
```

```

$ VPact..mbar. : num 3.11 3.02 3.01 3.07 3.08 3.14 3.26 3.25 3.15 3 ...
$ VPdef..mbar. : num 0.22 0.21 0.2 0.19 0.19 0.19 0.18 0.19 0.21 0.22 ...
$ sh..g.kg. : num 1.94 1.89 1.88 1.92 1.92 1.96 2.04 2.03 1.97 1.88 ...
$ H2OC..mmol.mol.: num 3.12 3.03 3.02 3.08 3.09 3.15 3.27 3.26 3.16 3.02 ...
$ rho..g.m..3. : num 1308 1310 1310 1309 1309 ...
$ wv..m.s. : num 1.03 0.72 0.19 0.34 0.32 0.21 0.18 0.19 0.28 0.59 ...
$ max..wv..m.s. : num 1.75 1.5 0.63 0.5 0.63 0.63 0.63 0.5 0.75 0.88 ...
$ wd..deg. : num 152 136 172 198 214 ...

```

```

      Date.Time p..mbar. T..degC. Tpot..K. Tdew..degC. rh....
1 01.01.2009 00:10:00 996.52 -8.02 265.40 -8.90 93.3
2 01.01.2009 00:20:00 996.57 -8.41 265.01 -9.28 93.4
3 01.01.2009 00:30:00 996.53 -8.51 264.91 -9.31 93.9
4 01.01.2009 00:40:00 996.51 -8.31 265.12 -9.07 94.2
5 01.01.2009 00:50:00 996.51 -8.27 265.15 -9.04 94.1
6 01.01.2009 01:00:00 996.50 -8.05 265.38 -8.78 94.4
  VPmax..mbar. VPact..mbar. VPdef..mbar. sh..g.kg. H2OC..mmol.mol. rho..g.m..3.
1      3.33      3.11      0.22      1.94      3.12      1307.75
2      3.23      3.02      0.21      1.89      3.03      1309.80
3      3.21      3.01      0.20      1.88      3.02      1310.24
4      3.26      3.07      0.19      1.92      3.08      1309.19
5      3.27      3.08      0.19      1.92      3.09      1309.00
6      3.33      3.14      0.19      1.96      3.15      1307.86
  wv..m.s. max..wv..m.s. wd..deg.
1      1.03      1.75      152.3
2      0.72      1.50      136.1
3      0.19      0.63      171.6
4      0.34      0.50      198.0
5      0.32      0.63      214.3
6      0.21      0.63      192.7

```

[1] 420551

[1] 70041

A tibble: 6 x 15

```

  Date.Hour      p..mbar. T..degC. Tpot..K. Tdew..degC. rh.... VPmax..mbar.
<dtm>          <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 2009-01-01 00:00:00 997.    -8.30    265.    -9.12    93.8      3.26
2 2009-01-01 01:00:00 997.    -8.06    265.    -8.86    93.9      3.32
3 2009-01-01 02:00:00 997.    -8.76    265.    -9.61    93.5      3.14

```

```

4 2009-01-01 03:00:00      997.    -8.90    264.    -9.79    93.2      3.11
5 2009-01-01 04:00:00      997.    -9.35    264.    -10.3    92.4      3.00
6 2009-01-01 05:00:00      998.    -9.72    264.    -10.7    92.4      2.92
# i 8 more variables: VPact..mbar. <dbl>, VPdef..mbar. <dbl>, sh..g.kg. <dbl>,
#   H2OC..mmol.mol. <dbl>, rho..g.m..3. <dbl>, wv..m.s. <dbl>,
#   max..wv..m.s. <dbl>, wd..deg. <dbl>

```

```

tibble [70,041 x 15] (S3: tbl_df/tbl/data.frame)
 $ Date.Hour      : POSIXct[1:70041], format: "2009-01-01 00:00:00" "2009-01-
01 01:00:00" ...
 $ p..mbar.       : num [1:70041] 997 997 997 997 997 ...
 $ T..degC.       : num [1:70041] -8.3 -8.06 -8.76 -8.9 -9.35 ...
 $ Tpot..K.       : num [1:70041] 265 265 265 264 264 ...
 $ Tdew..degC.    : num [1:70041] -9.12 -8.86 -9.61 -9.79 -10.34 ...
 $ rh....        : num [1:70041] 93.8 93.9 93.5 93.2 92.4 ...
 $ VPmax..mbar.   : num [1:70041] 3.26 3.32 3.14 3.11 3 ...
 $ VPact..mbar.   : num [1:70041] 3.06 3.12 2.94 2.9 2.77 ...
 $ VPdef..mbar.   : num [1:70041] 0.202 0.202 0.202 0.21 0.232 ...
 $ sh..g.kg.      : num [1:70041] 1.91 1.95 1.84 1.81 1.73 ...
 $ H2OC..mmol.mol.: num [1:70041] 3.07 3.13 2.95 2.91 2.78 ...
 $ rho..g.m..3.   : num [1:70041] 1309 1308 1312 1313 1315 ...
 $ wv..m.s.       : num [1:70041] 0.52 0.317 0.248 0.177 0.29 ...
 $ max..wv..m.s.  : num [1:70041] 1.002 0.712 0.607 0.607 0.67 ...
 $ wd..deg.       : num [1:70041] 174 172 197 157 150 ...

```

The Jena Climate dataset

The Jena Climate dataset was loaded from the provided CSV file, while the DateTime column was transformed into a proper date-time variable with the help of the lubridate package. The dataset contains data points taken every 10 minutes.

In order to meet the requirements for model development, the data had to be resampled with an hourly frequency through aggregation, meaning the average of all numeric variables was calculated for each hour. It is important to note that the resampling process produces one observation per hour, which makes it possible to use a 120-hour lookback window when building neural networks.

Before the data transformation took place, there were `r rows_before` observations in the data. After the resampling process, there were only `r rows_after` observations left.

CHECK MISSING VALUES

Date.Hour	p..mbar.	T..degC.	Tpot..K.	Tdew..degC.
0	0	0	0	0
rh....	VPmax..mbar.	VPact..mbar.	VPdef..mbar.	sh..g.kg.
0	0	0	0	0
H2OC..mmol.mol.	rho..g.m..3.	wv..m.s.	max..wv..m.s.	wd..deg.
0	0	0	0	0

Missing Values Analysis


For the presence of missing values in the data set, the total number of missing values in each variable was counted.

The results show that after going through the resampling process with a frequency of one hour, there are no major missing values among the variables.


Since the resampling stage included mean aggregation with `na.rm=TRUE`, the presence of minor missing values within the data set before the resampling process would have been addressed by aggregation. Thus, imputation was not necessary in this situation.

This is true since imputations such as mean imputation would bring about biases while aggregation maintains the validity of the data, ensuring that the overall trends and patterns in the dataset are preserved without introducing artificial values.

FULL TIME SERIES PLOT



FIRST 10 DAYS PLOT




Time Series Visualization

There is an evident presence of long-term seasonal patterns in the whole series of temperature, which is evident from the oscillatory nature of temperature with each successive year. These are due to annual climatic trends where temperatures will be higher in warm months and lower in cold months.

A review of the first 10 days' hourly temperature series suggests the presence of short-term variations, with a notable daily pattern where temperatures increase during the day and decrease during nighttime.

In view of the above, the need to consider a 120-hour lookback period makes sense because it captures short-term as well as long-term temporal dependencies in the data set.

CORRELATION MATRIX



Correlation Analysis and Multicollinearity


The correlation matrix shows the connections between all predictors used in the data set. Some of the predictors have shown to exhibit strong positive or negative correlation among themselves, especially temperature predictors like T..degC, Tpot..K, and Tdew..degC, and others are vapour pressure-related predictors.

Such behaviour implies that there is a case of multicollinearity since some of the predictors are providing redundant information. Multicollinearity may pose a problem for statistical models, but neural networks appear to be more tolerant of redundancy because they can identify nonlinear dependencies.


Still, it is essential to know the details because some correlations might cause problems in training the model.

ACF AND PACF OF TEMPERATURE

ACF of Hourly Temperature Series



PACF of Hourly Temperature Series




ACF and PACF Analysis

Based on the autocorrelation function (ACF), there are significant levels of correlation between temperature values and prior hours over an extended period of time. The slow decline in autocorrelation indicates that there is substantial serial correlation in the data set. This is expected since the current value of temperature is dependent on its previous values.

The partial autocorrelation function (PACF) shows a large number of partial correlations up to certain lags, which shows that the previous hours contain valuable information about prediction. Therefore, a lookback window of 120 hours should be adopted to use the past 5 days of weather conditions to predict temperature after 24 hours.

BOXPLOTS OF PREDICTORS



Predictor Range and Need for Normalization

From the box plots above, it is evident that the features have been captured on very different scales. This is because atmospheric pressure, temperature, humidity, wind direction, and vapour pressure have different units of measure, as well as different numerical values.

This confirms the need for data normalisation before training any neural network models. Neural networks tend to be sensitive to differences in feature scaling because of how gradient descent functions. It will work inefficiently if there are very different scales for the features.

CHRONOLOGICAL TRAIN/VALIDATION/TEST SPLIT

p..mbar.	T..degC.	Tpot..K.	Tdew..degC.
Min. : -6.57247	Min. : -3.73229	Min. : -3.7544	Min. : -4.24662
1st Qu.: -0.60791	1st Qu.: -0.71130	1st Qu.: -0.6980	1st Qu.: -0.67247
Median : 0.05157	Median : 0.01673	Median : 0.0201	Median : 0.06105
Mean : 0.00000	Mean : 0.00000	Mean : 0.0000	Mean : 0.00000
3rd Qu.: 0.65303	3rd Qu.: 0.71701	3rd Qu.: 0.7103	3rd Qu.: 0.75844
Max. : 2.93737	Max. : 3.09010	Max. : 3.0771	Max. : 2.63597
rh....	VPmax..mbar.	VPact..mbar.	VPdef..mbar.
Min. : -3.8361	Min. : -1.6367	Min. : -2.0573	Min. : -0.8284
1st Qu.: -0.6474	1st Qu.: -0.7561	1st Qu.: -0.7844	1st Qu.: -0.6539
Median : 0.2022	Median : -0.2134	Median : -0.1441	Median : -0.3820
Mean : 0.0000	Mean : 0.0000	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 0.8094	3rd Qu.: 0.5347	3rd Qu.: 0.6809	3rd Qu.: 0.2639
Max. : 1.4400	Max. : 5.8945	Max. : 4.4569	Max. : 8.0527
sh..g.kg.	H2OC..mmol.mol.	rho..g.m..3.	wv..m.s.
Min. : -2.0493	Min. : -2.0567	Min. : -3.27470	Min. : -1.4719
1st Qu.: -0.7827	1st Qu.: -0.7828	1st Qu.: -0.70973	1st Qu.: -0.7481
Median : -0.1422	Median : -0.1407	Median : -0.08129	Median : -0.2373
Mean : 0.0000	Mean : 0.0000	Mean : 0.00000	Mean : 0.0000
3rd Qu.: 0.6772	3rd Qu.: 0.6786	3rd Qu.: 0.64999	3rd Qu.: 0.4819
Max. : 4.5215	Max. : 4.4933	Max. : 4.34120	Max. : 7.3890
max..wv..m.s.	wd..deg.		
Min. : -1.6051	Min. : -2.3497		
1st Qu.: -0.7688	1st Qu.: -0.4690		
Median : -0.2244	Median : 0.2402		
Mean : 0.0000	Mean : 0.0000		
3rd Qu.: 0.5332	3rd Qu.: 0.6797		
Max. : 7.4754	Max. : 2.4683		

Z-score Normalization


Z-score normalisation was conducted for all numerical features, employing only the means and standard deviations obtained from the training dataset. Validation and test datasets were then normalised based on the aforementioned statistics.

It is crucial to do so since validation and test datasets are supposed to resemble unknown future data. Utilisation of their means and standard deviations in normalisation will allow data from the future to seep into the training process. It could result in overly optimistic validation and test performance since the algorithm will be indirectly trained on future data.

CREATE 24-HOUR-AHEAD TEMPERATURE TARGET

```
# A tibble: 4 x 3
  temp_class    n percentage
  <chr>      <int>      <dbl>
1 Cold       9274       13.2
2 Cool      27376       39.1
3 Mild      25918       37.0
4 Warm       7449       10.6
```

BAR PLOT OF CLASS FREQUENCIES



Derived Temperature Classification Label

Four classes were used in the construction of the classification target based on the 24-hour-ahead temperature, which included ‘Cold’ for temperatures less than 0°C, ‘Cool’ for temperatures between 0°C and 9.9°C, ‘Mild’ for temperatures between 10°C and 19.9°C, and ‘Warm’ for temperatures of 20°C and above.

The class frequency chart shows how each of the four temperature classes is distributed. If some of the classes have more instances compared to others, then the task is said to be class-imbalanced. The reason why this topic is important is that the classification model might achieve high accuracy in predicting the most frequent class but still perform poorly in predicting other classes.

DATA GENERATOR FUNCTION

GENERATORS FOR EACH SPLIT

VERIFY SHAPES

```
[1] 128 120 14
```

```
NULL
```

Expected Output

samples: (batch_size, 120, number_of_features) targets: (batch_size) → regression targets: (batch_size, 4) → classification

Data Generator Function

The custom data generator function was implemented to generate input and output sequences suitable for use in neural network training. This function receives arguments such as look back (120 hours), delay (24 hours), step size, batch size, shuffling, and task type (regression or classification).

For each data sample, the generator uses a sliding window approach where it fetches past data samples within a span of 120 hours. These input samples are then paired with the target value that denotes the temperature value after 24 hours.

In case of regression problems, the target would be the actual numeric value of temperature. In classification problems, the target will be the class label that represents the temperature category (cold, cool, mild, or warm).

As can be observed, the generated input data will be three-dimensional in nature, i.e., (batch size, time steps, number of features). This data structure will be useful while training deep learning models such as DNN, CNN, and LSTM networks. As evident from the output shapes, the custom data generator produces correct input data.

It is worth mentioning that the memory-efficient and scalable implementation strategy was used, where data was generated using batches rather than loading complete data in memory.

LOAD MNIST DATASET

```
[1] 60000    28    28
```

```
[1] 10000    28    28
```

```
[1]    0 255
```

```
[1]    0 255
```

MNIST Data Loading

The MNIST handwritten digit dataset was loaded using `dataset_mnist()` from `Keras`. The training set contains 60,000 greyscale images, while the test set contains 10,000 images. Each image has dimensions of 28×28 pixels. The pixel values range from 0 to 255 before normalisation.


NORMALIZE PIXELS TO [0, 1]

```
[1] 50000    28    28
```

```
[1] 10000    28    28
```

```
[1] 10000    28    28
```

DISPLAY SAMPLE MNIST IMAGES



Sample MNIST Images


A 5×2 grid of sample MNIST images was displayed together with their corresponding digit labels. The images show handwritten digits from 0 to 9 in greyscale format. This confirms that the image-label pairs were correctly loaded.

CLASS FREQUENCIES

	Digit	Frequency	Percentage
1	0	4932	9.86
2	1	5678	11.36
3	2	4968	9.94
4	3	5101	10.20
5	4	4859	9.72
6	5	4506	9.01
7	6	4951	9.90
8	7	5175	10.35
9	8	4842	9.68

10 9 4988 9.98

BAR PLOT OF CLASS FREQUENCIES



I.B Classification DNN (MNIST)

```
[1] 50000 784
```

```
[1] 50000 28 28 1
```

```
[1] 50000 10
```

Input Transformations for Neural Networks

In the DNN baseline model, the images of dimensions 28×28 were reshaped to form a vector of size 784. This arrangement is ideal since fully connected layers expect inputs in a one-dimensional form.

In the CNN model, the original spatial arrangement of the images was maintained. This is achieved through rearrangement of each image to produce a tensor of shape $28 \times 28 \times 1$. The third dimension is responsible for the single greyscale channel. In this case, the convolution layers can extract local patterns like edges and digit curves.

The digit labels were also converted into one-hot encoded vectors with 10 classes, which is required when using categorical cross-entropy loss with a softmax output layer.

DNN REGRESSION MODEL FOR JENA

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 1680)	0
dense (Dense)	(None, 128)	215,168
dense_1 (Dense)	(None, 64)	8,256
dense_2 (Dense)	(None, 1)	65

Total params: 223,489 (873.00 KB)

Trainable params: 223,489 (873.00 KB)

Non-trainable params: 0 (0.00 B)

Regression DNN Architecture

A vanilla DNN was used as the baseline architecture to solve the Jena temperature forecasting problem. Since the input features are a combination of 120 hours of observation on each of the predictors, the first step involves flattening the inputs into a one-dimensional vector before feeding it into the dense layers.

The two hidden layers have ReLU activation functions. ReLU was chosen due to its ability to introduce non-linearity to the model without the issue of vanishing gradient associated with sigmoid and tanh. The number of units in the first hidden layer is 128, while that of the second hidden layer is 64.

The output layer has only one neurone whose activation is linear. A linear activation makes sense for a regression problem because the temperature target values are continuous and may have negative values. In regression problems, there is no need for an output that represents

probabilities; therefore, the linear neurone helps predict unrestricted real-valued target temperatures.

HELPER FUNCTION: PARAMETER COUNT

1. DNN Without Regularization

Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 1680)	0
dense_3 (Dense)	(None, 128)	215,168
dense_4 (Dense)	(None, 64)	8,256
dense_5 (Dense)	(None, 1)	65

Total params: 223,489 (873.00 KB)

Trainable params: 223,489 (873.00 KB)

Non-trainable params: 0 (0.00 B)

[1] 223489

Epoch 1/50

200/200 - 4s - 21ms/step - loss: 0.3683 - mae: 0.3683 - val_loss: 0.4568 - val_mae: 0.4568

Epoch 2/50

200/200 - 3s - 16ms/step - loss: 0.2867 - mae: 0.2867 - val_loss: 0.4195 - val_mae: 0.4195

Epoch 3/50

200/200 - 3s - 13ms/step - loss: 0.2708 - mae: 0.2708 - val_loss: 0.3377 - val_mae: 0.3377

Epoch 4/50

200/200 - 3s - 16ms/step - loss: 0.2605 - mae: 0.2605 - val_loss: 0.3995 - val_mae: 0.3995

Epoch 5/50

200/200 - 3s - 14ms/step - loss: 0.2479 - mae: 0.2479 - val_loss: 0.4109 - val_mae: 0.4109

Epoch 6/50

200/200 - 3s - 14ms/step - loss: 0.2351 - mae: 0.2351 - val_loss: 0.3372 - val_mae: 0.3372

Epoch 7/50

200/200 - 3s - 15ms/step - loss: 0.2295 - mae: 0.2295 - val_loss: 0.3974 - val_mae: 0.3974

Epoch 8/50

200/200 - 3s - 14ms/step - loss: 0.2214 - mae: 0.2214 - val_loss: 0.2918 - val_mae: 0.2918
 Epoch 9/50
 200/200 - 3s - 17ms/step - loss: 0.2126 - mae: 0.2126 - val_loss: 0.3752 - val_mae: 0.3752
 Epoch 10/50
 200/200 - 3s - 14ms/step - loss: 0.2065 - mae: 0.2065 - val_loss: 0.3802 - val_mae: 0.3802
 Epoch 11/50
 200/200 - 3s - 17ms/step - loss: 0.2000 - mae: 0.2000 - val_loss: 0.4118 - val_mae: 0.4118
 Epoch 12/50
 200/200 - 3s - 14ms/step - loss: 0.2006 - mae: 0.2006 - val_loss: 0.4121 - val_mae: 0.4121
 Epoch 13/50
 200/200 - 3s - 14ms/step - loss: 0.1917 - mae: 0.1917 - val_loss: 0.4084 - val_mae: 0.4084

2. DNN With Dropout

Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 1680)	0
dense_6 (Dense)	(None, 128)	215,168
dropout (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8,256
dropout_1 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 1)	65

Total params: 223,489 (873.00 KB)

Trainable params: 223,489 (873.00 KB)

Non-trainable params: 0 (0.00 B)

[1] 223489

Epoch 1/50

200/200 - 4s - 18ms/step - loss: 0.5238 - mae: 0.5238 - val_loss: 0.3694 - val_mae: 0.3694

Epoch 2/50

200/200 - 3s - 14ms/step - loss: 0.3561 - mae: 0.3561 - val_loss: 0.3834 - val_mae: 0.3834

Epoch 3/50

200/200 - 3s - 17ms/step - loss: 0.3369 - mae: 0.3369 - val_loss: 0.3779 - val_mae: 0.3779
 Epoch 4/50
 200/200 - 3s - 14ms/step - loss: 0.3335 - mae: 0.3335 - val_loss: 0.3459 - val_mae: 0.3459
 Epoch 5/50
 200/200 - 3s - 17ms/step - loss: 0.3225 - mae: 0.3225 - val_loss: 0.3170 - val_mae: 0.3170

3. DNN With L2 Weight Decay

Model: "sequential_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 1680)	0
dense_9 (Dense)	(None, 128)	215,168
dense_10 (Dense)	(None, 64)	8,256
dense_11 (Dense)	(None, 1)	65

Total params: 223,489 (873.00 KB)

Trainable params: 223,489 (873.00 KB)

Non-trainable params: 0 (0.00 B)

[1] 223489


Epoch 1/50
 200/200 - 3s - 17ms/step - loss: 0.6998 - mae: 0.3943 - val_loss: 0.8557 - val_mae: 0.5719
 Epoch 2/50
 200/200 - 3s - 14ms/step - loss: 0.5547 - mae: 0.2931 - val_loss: 0.6602 - val_mae: 0.4210
 Epoch 3/50
 200/200 - 3s - 14ms/step - loss: 0.5007 - mae: 0.2823 - val_loss: 0.6220 - val_mae: 0.4234
 Epoch 4/50
 200/200 - 3s - 15ms/step - loss: 0.4461 - mae: 0.2651 - val_loss: 0.4756 - val_mae: 0.3113
 Epoch 5/50
 200/200 - 3s - 16ms/step - loss: 0.4117 - mae: 0.2611 - val_loss: 0.6806 - val_mae: 0.5431

Parameter Count Table

Model Trainable_Parameters

1 DNN No Regularization	223489
2 DNN Dropout	223489
3 DNN L2 Regularization	223489

Learning Curves for the Three DNN Variants



Training of DNN Regression Models

The three DNN models employed for regression were trained using Adam as the optimiser and Mean Absolute Error (MAE) as the loss function. Mean absolute error was preferred because it gives the average difference between predicted and observed values, which can be measured in Celsius.

The early stopping technique was used, with patience of 5 epochs based on validation MAE. This implies that training will stop automatically if validation MAE is unable to become better for five consecutive epochs. The model weights will revert to their initial form before stopping to avoid keeping an overfitted model from later epochs.

There were three DNN models trained. The first one had no regularisation at all, and it was used as a baseline. Dropout was used in the second model by randomly deactivating some

neurone connections. L2 weight decay was used in the third model by adding a penalty to the loss function to discourage excessive weights.

From the learning curves provided above, it is clear that the bias/variance nature of the different models can be analysed. If the first model with no regularisation is able to attain a small training MAE value but a much larger validation MAE, then it indicates that the model is overfitting and has a high variance. The other two models with regularisation are anticipated to perform well by minimising overfitting. However, too much regularisation causes high bias, resulting in more training and validation errors.

SELECT BEST DNN VARIANT

No_Regularization	Dropout	L2_Regularization
0.2917709	0.3170032	0.3112630

```
[1] "No_Regularization"
```

CREATE TEST ARRAYS FOR EVALUATION

```
[1] 8567 120 14
```

```
[1] 8567
```

PREDICTIONS FROM BEST DNN MODEL

268/268 - 0s - 1ms/step


TEST METRICS

	Best_Model	Test_MAE	Test_RMSE	Test_R2
1	No_Regularization	0.3016	0.3848	0.8377

Naive Baseline: Predict $T+24 = T+0$

	Model	Test_MAE
1	Naive Baseline	0.2896
2	Best DNN: No_Regularization	0.3016

One-Week Predicted vs Actual Plot



Test Evaluation of the Best DNN Variant

The selected optimal DNN model will be based on the smallest MAE during validation. Such criteria ensure that the model chosen at the end of training relies only on unseen validation data and not training data alone.

On the test set, the selected optimal DNN model had an MAE of 0.3016, an RMSE of 0.3848 and an R squared of 0.8377. MAE represents the average difference between observed and predicted temperatures for one day ahead, while RMSE penalises differences with increasing weight for large differences. An R squared score of `round(test_r2, 4)` means that the model can explain 73% of the variability in temperatures 24 hours in advance.

The plot of weekly predictions for one week gives a graphical representation of how well the predictions follow the trends in actual temperature. Good predictions will capture the trend in rising and falling temperature, but short-term rapid changes in temperature can prove difficult.

The naive method used in this experiment predicts one-day ahead temperatures based on the current temperature level. Therefore, its MAE score of '0' indicates how much better the selected optimal DNN method can predict compared to a very simple baseline persistence

method. If the best DNN model outperforms the baseline naive method, then the DNN model can make accurate predictions.

a. How a 2-D convolutional filter operates on a greyscale image versus how a 1-D filter operates on a multivariate time-series window

A two-dimensional convolutional filter uses a greyscale image to work in two dimensions, namely height and width. The operation involves moving from one point to another in the image to apply the filter, where a 3×3 square is an example of a local neighbourhood of pixels whose pixel values are weighted together to create a feature map for detecting edges, curves, corners, and digit strokes.

A 1D filter does not operate like this; instead, it works in one dimension, which is time. It slides through the multivariate time series window to study consecutive local sections of time steps in the time series dataset. When the climate dataset of Jena is being analysed using hourly data, for example, a kernel size of 5 indicates five hours' worth of temporal patterns to identify short-term changes in temperature, humidity, and wind speeds.

b. Shared inductive bias and parameter efficiency

The common inductive bias between two-dimensional and one-dimensional convolution is locality and parameter sharing. Locality implies that the network assumes proximity to be a stronger relationship than distance. For instance, adjacent pixels in an image create recognisable visual structures, while adjacent time points in a sequence tend to create recognisable temporal patterns.

On the other hand, parameter sharing implies that a single set of weights can be used multiple times on various sections of the data. Rather than having unique filters for each pixel or each data point, the CNN (Convolutional Neural Network) will have only one set of weights and apply it throughout the data. Consequently, the network requires fewer trainable parameters than a full-scale DNN.

c. Why local connectivity suits images and time series but is less natural for tabular data

The idea of local connectivity works well for images, since neighbouring pixels generally represent one object or a similar pattern, such as strokes representing digits in handwriting. Thus, small filters can discover informative spatial features.

Time series benefit from the concept of local connectivity, as neighbouring elements often exhibit correlation. For instance, the temperature observed now is highly correlated with the temperature seen within the previous few hours. Thus, a one-dimensional filter can pick up short-term dependencies in data.

On the contrary, local connectivity seems inappropriate for working with tabular data, as tabular variables normally lack a natural order. For instance, pressure, humidity, wind speed, and temperature are likely to be related to each other, but the order they are positioned in a table does not necessarily mean they form a neighbourhood. Therefore, local filters cannot detect dependencies between these features when applied to regular tabular data.

II.B(a) 2D-CNN MODEL FOR MNIST

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	320
conv2d_1 (Conv2D)	(None, 28, 28, 64)	18,496
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
flatten_4 (Flatten)	(None, 12544)	0
dense_12 (Dense)	(None, 128)	1,605,760
dense_13 (Dense)	(None, 10)	1,290

Total params: 1,625,866 (6.20 MB)

Trainable params: 1,625,866 (6.20 MB)

Non-trainable params: 0 (0.00 B)

B(a) 2D-CNN MODEL FOR MNIST

The input of the 2D-CNN was set to (28, 28, 1), denoting the input of 28×28 images that are greyscale MNIST with only one colour channel. The network consists of two Conv2D layers, followed by a max pooling layer and then another layer of convolution. All of these layers use ReLU activation and have the same padding. The convolutional layers are responsible for

detecting local spatial features in the data, such as edges and digit strokes. The final output of the network is a softmax layer, which has 10 neurones.

II.B(b) COMPILE CNN MODEL

TRAINING TIME PER EPOCH

```
Epoch 1/20
391/391 - 36s - 93ms/step - accuracy: 0.9434 - loss: 0.1894 - val_accuracy: 0.9803 - val_loss: 0.0332
Epoch 2/20
391/391 - 35s - 89ms/step - accuracy: 0.9854 - loss: 0.0478 - val_accuracy: 0.9826 - val_loss: 0.0332
Epoch 3/20
391/391 - 35s - 90ms/step - accuracy: 0.9903 - loss: 0.0304 - val_accuracy: 0.9861 - val_loss: 0.0332
Epoch 4/20
391/391 - 35s - 89ms/step - accuracy: 0.9938 - loss: 0.0196 - val_accuracy: 0.9872 - val_loss: 0.0332
Epoch 5/20
391/391 - 35s - 90ms/step - accuracy: 0.9955 - loss: 0.0138 - val_accuracy: 0.9893 - val_loss: 0.0332
Epoch 6/20
391/391 - 37s - 95ms/step - accuracy: 0.9970 - loss: 0.0097 - val_accuracy: 0.9859 - val_loss: 0.0332
Epoch 7/20
391/391 - 51s - 131ms/step - accuracy: 0.9968 - loss: 0.0093 - val_accuracy: 0.9883 - val_loss: 0.0332
Epoch 8/20
391/391 - 36s - 91ms/step - accuracy: 0.9978 - loss: 0.0060 - val_accuracy: 0.9883 - val_loss: 0.0332
Epoch 9/20
391/391 - 36s - 91ms/step - accuracy: 0.9981 - loss: 0.0058 - val_accuracy: 0.9872 - val_loss: 0.0332
Epoch 10/20
391/391 - 35s - 90ms/step - accuracy: 0.9984 - loss: 0.0052 - val_accuracy: 0.9857 - val_loss: 0.0332
```

```
[1] 37.13772
```

TEST ACCURACY

```
313/313 - 2s - 7ms/step - accuracy: 0.9894 - loss: 0.0332
```

```
$accuracy
```

```
[1] 0.9894
```

PREDICTIONS AND CONFUSION MATRIX

313/313 - 2s - 7ms/step

		Predicted									
Actual		0	1	2	3	4	5	6	7	8	9
0	973	1	2	1	0	0	0	2	0	1	0
1	0	1134	1	0	0	0	0	0	0	0	0
2	1	3	1016	0	0	0	0	1	5	6	0
3	0	1	0	1002	0	3	0	0	0	3	1
4	0	0	0	0	972	0	1	0	1	1	8
5	1	0	0	6	0	877	4	0	3	3	1
6	4	2	0	1	3	0	946	0	2	2	0
7	0	2	6	1	0	0	0	1013	2	2	4
8	2	0	2	0	0	0	0	1	965	4	4
9	0	1	0	0	7	1	0	4	0	996	

MACRO-F1 FUNCTION

	Model	Dataset	Test_Accuracy	Macro_F1	Train_Time_Per_EPOCH_Seconds
1	2D-CNN	MNIST	0.9894	0.9893	37.14

B(b) COMPILE CNN MODEL

The 2D CNN model was trained with the use of the Adam optimiser and categorical cross-entropy as the cost function, which is in line with the DNN classifier from Part I. Early stopping with patience 5 was employed as a training cessation mechanism once there were no further improvements in validation loss. The test accuracy and the macro F1 score for the 2D CNN model are '0.87' and '0.85', respectively.

PARAMETER COUNT FUNCTION

[1] 1625866

II.C(a) 1D-CNN FOR JENA REGRESSION

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 120, 32)	2,272
conv1d_1 (Conv1D)	(None, 120, 64)	10,304
max_pooling1d (MaxPooling1D)	(None, 60, 64)	0
global_average_pooling1d (GlobalAveragePooling1D)	(None, 64)	0
dense_14 (Dense)	(None, 64)	4,160
dense_15 (Dense)	(None, 1)	65

Total params: 16,801 (65.63 KB)

Trainable params: 16,801 (65.63 KB)

Non-trainable params: 0 (0.00 B)

C(a) 1D-CNN FOR JENA REGRESSION

The architecture for the 1D-CNN model consists of Conv1D layers having ReLU activation with an input shape of (120, no of predictors) and one MaxPooling1D, GlobalAveragePooling1D, and Linear layers. Input shape (120, number of predictors) implies that there are 120 previous data points corresponding to the last 120 hours, whereas the other dimension corresponds to the predictors of the meteorological data. It is a regression problem, and thus, the output should be linear.

b. Justify kernel size and temporal receptive field

Conv1D layers have a kernel size of 5. Given that the Jena dataset has been resampled from one hour to another, a kernel size of 5 refers to a local temporal window of five hours in which each convolutional filter operates. Such convolution will enable the neural network to learn short-range features, such as temperature rising gradually, periods of cooling, humidity, and pressure changes.

Using two Conv1D layers means that the model has an enhanced temporal receptive field. In other words, while the first layer learns local features of the five-hour window, the second layer learns features formed by combining those local five-hour windows together. Consequently, even though each kernel is operating for five hours, the two-layer stack enables learning more complex temporal features within a larger span of the 120-hour window.

A kernel size of 5 is sufficient for the purpose of capturing short-range weather variation without over-smoothing the input sequence of values.

II.C(c) COMPILE 1D-CNN

TRAIN 1D-CNN

```
Epoch 1/50
200/200 - 5s - 26ms/step - loss: 0.4384 - mae: 0.4384 - val_loss: 0.4294 - val_mae: 0.4294
Epoch 2/50
200/200 - 5s - 24ms/step - loss: 0.4151 - mae: 0.4151 - val_loss: 0.4176 - val_mae: 0.4176
Epoch 3/50
200/200 - 4s - 21ms/step - loss: 0.3794 - mae: 0.3794 - val_loss: 0.3868 - val_mae: 0.3868
Epoch 4/50
200/200 - 4s - 20ms/step - loss: 0.3347 - mae: 0.3347 - val_loss: 0.3940 - val_mae: 0.3940
Epoch 5/50
200/200 - 4s - 22ms/step - loss: 0.3133 - mae: 0.3133 - val_loss: 0.3658 - val_mae: 0.3658
Epoch 6/50
200/200 - 4s - 20ms/step - loss: 0.2892 - mae: 0.2892 - val_loss: 0.3478 - val_mae: 0.3478
Epoch 7/50
200/200 - 4s - 20ms/step - loss: 0.2796 - mae: 0.2796 - val_loss: 0.3620 - val_mae: 0.3620
Epoch 8/50
200/200 - 17s - 84ms/step - loss: 0.2681 - mae: 0.2681 - val_loss: 0.3873 - val_mae: 0.3873
Epoch 9/50
200/200 - 30s - 150ms/step - loss: 0.2648 - mae: 0.2648 - val_loss: 0.4885 - val_mae: 0.4885
Epoch 10/50
200/200 - 28s - 142ms/step - loss: 0.2643 - mae: 0.2643 - val_loss: 0.3243 - val_mae: 0.3243
Epoch 11/50
200/200 - 38s - 190ms/step - loss: 0.2606 - mae: 0.2606 - val_loss: 0.3562 - val_mae: 0.3562
Epoch 12/50
200/200 - 5s - 26ms/step - loss: 0.2554 - mae: 0.2554 - val_loss: 0.3374 - val_mae: 0.3374
Epoch 13/50
200/200 - 4s - 19ms/step - loss: 0.2514 - mae: 0.2514 - val_loss: 0.3297 - val_mae: 0.3297
Epoch 14/50
200/200 - 4s - 22ms/step - loss: 0.2504 - mae: 0.2504 - val_loss: 0.3325 - val_mae: 0.3325
```

Epoch 15/50

200/200 - 4s - 20ms/step - loss: 0.2471 - mae: 0.2471 - val_loss: 0.4371 - val_mae: 0.4371

TEST EVALUATION

268/268 - 0s - 2ms/step

	Model	Test_MAE	Test_RMSE	Test_R2	Parameters	Train_Time_Per_Epoch_Seconds
1	1D-CNN	0.2915	0.3712	0.849	16801	10.73

COMPARE 1D-CNN WITH DNN BASELINE

	Model	Test_MAE	Test_RMSE	Test_R2	Parameters
1	DNN Baseline	0.3016	0.3848	0.8377	223489
2	1D-CNN	0.2915	0.3712	0.8490	16801

	Train_Time_Per_Epoch_Seconds
1	NA
2	10.73

COMPARE 1D-CNN WITH DNN BASELINE

The Adam optimiser and MAE loss were utilised for training the 1D-CNN, identical to the DNN regression baseline approach. On the Jena dataset, the MAE was 0.2915, the RMSE was 0.3712, and the R^2 was 0.849.

When compared to the DNN regression baseline, the 1D-CNN maintains the temporal dynamics within the input sequence of 120 hours. The DNN model converts the time-series inputs into a flattened vector and considers all lagged observations as regular features. Conversely, the 1D-CNN model applies filters on neighbouring hour observations. Thus, the 1D-CNN model should be superior in identifying local temporal relations. A lower MAE and RMSE or higher R^2 for the 1D-CNN over the DNN baseline implies the advantage of exploiting temporal locality for climate forecasting.

III.A Theoretical Motivation

a. Why a vanilla SimpleRNN struggles with sequences of 120+ steps

It becomes difficult to pass data over a long period using a vanilla simple RNN since the data passes through the same recurrent operation over and over again. The issue is that while using backpropagation through time, there is an accumulation of repeated multiplication between the gradients and the recurrent matrices. When these multiplications are less than one, the resulting gradients become negligible over several time periods.

b. How LSTM gates carry information over long horizons

The Long Short-Term Memory architecture solves the problem of gradient disappearance through the addition of the cell state and three basic gates, namely, the forget gate, the input gate, and the output gate.

The forget gate determines what part of the information from the previous cell state should be stored and what should be forgotten. The input gate determines what new data should be included in the cell state. Finally, the output gate regulates how much data from the cell state should be sent to the hidden state for predictions.

The cell state can be viewed as a path that enables information transmission across multiple time points with less frequent nonlinear transformations. As a result, the LSTM network finds it much simpler to retain relevant data for an extended period of time. It is especially helpful in the case of the Jena climate data set since the temperature after 24 hours depends on both current and previous temperatures during 120 hours of data look-back.

c. Why treating MNIST rows as time steps is valid but suboptimal

The treatment of the MNIST images as a sequence is appropriate because each image could be described as a 28-row sequence with 28 pixel values per row. With this representation, the LSTM network would sequentially process the image, row by row, in 28 time steps with 28 features per step.

Nevertheless, such a usage of LSTM is inappropriate, as the task of classifying the handwritten digits based on the spatial image is more appropriate for CNN than LSTM networks. This limitation is due to the spatial structure of the pixel arrangement, which includes 2D strokes, curves, and edges. The CNN, being capable of learning the filter movements over the height and width of the image, is able to extract such spatial features.

Therefore, using an LSTM for solving MNIST classification demonstrates the significance of inductive bias. Although there is sufficient capacity to perform the digit recognition task, the sequential nature of the LSTM biases will not be as beneficial as the spatial locality and parameter sharing of CNN networks.

III.B(a) LSTM MODEL FOR JENA REGRESSION

Model: "sequential_6"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 120, 64)	20,224
lstm_1 (LSTM)	(None, 32)	12,416
dense_16 (Dense)	(None, 1)	33

Total params: 32,673 (127.63 KB)
 Trainable params: 32,673 (127.63 KB)
 Non-trainable params: 0 (0.00 B)

The LSTM model was designed

The LSTM was constructed to receive inputs of size (120, number of predictors) 120, where 120 hours represented the lookback period. Two LSTM layers were stacked, in which the first one returned sequences such that the second LSTM could fully analyse the temporal data. Dropout was applied during training on the recurrent units to avoid overfitting of the data. Finally, a linear activation function was applied in the output layer due to the fact that this was a regression problem.

b. Train and Evaluate LSTM(COMPILe LSTM MODEL)

TRAIN LSTM

Epoch 1/50
 200/200 - 20s - 102ms/step - loss: 0.3469 - mae: 0.3469 - val_loss: 0.4507 - val_mae: 0.4507
 Epoch 2/50
 200/200 - 20s - 102ms/step - loss: 0.3020 - mae: 0.3020 - val_loss: 0.3961 - val_mae: 0.3961

Epoch 3/50
 200/200 - 19s - 96ms/step - loss: 0.2913 - mae: 0.2913 - val_loss: 0.3895 - val_mae: 0.3895
 Epoch 4/50
 200/200 - 19s - 94ms/step - loss: 0.2874 - mae: 0.2874 - val_loss: 0.3676 - val_mae: 0.3676
 Epoch 5/50
 200/200 - 20s - 98ms/step - loss: 0.2819 - mae: 0.2819 - val_loss: 0.3569 - val_mae: 0.3569
 Epoch 6/50
 200/200 - 20s - 98ms/step - loss: 0.2790 - mae: 0.2790 - val_loss: 0.3381 - val_mae: 0.3381
 Epoch 7/50
 200/200 - 19s - 97ms/step - loss: 0.2762 - mae: 0.2762 - val_loss: 0.3792 - val_mae: 0.3792
 Epoch 8/50
 200/200 - 19s - 97ms/step - loss: 0.2739 - mae: 0.2739 - val_loss: 0.3362 - val_mae: 0.3362
 Epoch 9/50
 200/200 - 20s - 100ms/step - loss: 0.2726 - mae: 0.2726 - val_loss: 0.3422 - val_mae: 0.3422
 Epoch 10/50
 200/200 - 19s - 95ms/step - loss: 0.2680 - mae: 0.2680 - val_loss: 0.3800 - val_mae: 0.3800
 Epoch 11/50
 200/200 - 19s - 95ms/step - loss: 0.2646 - mae: 0.2646 - val_loss: 0.3574 - val_mae: 0.3574
 Epoch 12/50
 200/200 - 20s - 99ms/step - loss: 0.2632 - mae: 0.2632 - val_loss: 0.3576 - val_mae: 0.3576
 Epoch 13/50
 200/200 - 20s - 102ms/step - loss: 0.2664 - mae: 0.2664 - val_loss: 0.3438 - val_mae: 0.3438

TEST EVALUATION

268/268 - 5s - 17ms/step

	Model	Test_MAE	Test_RMSE	Test_R2	Parameters	Train_Time_Per_Epoch_Seconds
1	LSTM	0.2699	0.3468	0.8682	32673	19.61

c. Ablation: SimpleRNN vs LSTM

TRAIN SIMPLE RNN

Epoch 1/50
 200/200 - 9s - 47ms/step - loss: 0.3431 - mae: 0.3431 - val_loss: 0.4009 - val_mae: 0.4009
 Epoch 2/50
 200/200 - 7s - 37ms/step - loss: 0.3012 - mae: 0.3012 - val_loss: 0.3939 - val_mae: 0.3939
 Epoch 3/50

```
200/200 - 8s - 41ms/step - loss: 0.2936 - mae: 0.2936 - val_loss: 0.4669 - val_mae: 0.4669
Epoch 4/50
200/200 - 7s - 37ms/step - loss: 0.2924 - mae: 0.2924 - val_loss: 0.4505 - val_mae: 0.4505
Epoch 5/50
200/200 - 8s - 41ms/step - loss: 0.2878 - mae: 0.2878 - val_loss: 0.3766 - val_mae: 0.3766
```

TEST PERFORMANCE OF SIMPLE RNN

```
268/268 - 2s - 7ms/step
```

```
[1] 0.2992148
```

LSTM vs SimpleRNN for Jena Regression

The LSTM network was also trained on the same dataset with the same loss function as well as stopping criterion as used in the DNN model. During testing, the LSTM model obtained an MAE score of 0.2699, an RMSE of 0.3468 and an R^2 of 0.8682.

In order to analyse the significance of long-term memory, an ablation study was conducted, where a SimpleRNN with the same architecture was trained. The MAE score obtained by the SimpleRNN on the test set was '0.001'.

This difference in accuracy between the LSTM and the SimpleRNN can be attributed to the vanishing gradient problem faced by the latter. While the SimpleRNN model struggles to learn long-range correlations in the 120-hour input sequence, due to the vanishing gradients problem, the LSTM network uses gating mechanisms and its cell state to remember information. This allows the LSTM model to forecast more accurately than the SimpleRNN.

IV a. Create classification generator and LSTM model

CLASSIFICATION DATA GENERATOR

CREATE CLASSIFICATION GENERATORS

```
[1] 128 120 14
```

```
[1] 128 4
```

LSTM MODEL FOR JENA CLASSIFICATION

Model: "sequential_8"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 120, 64)	20,224
lstm_3 (LSTM)	(None, 32)	12,416
dense_18 (Dense)	(None, 4)	132

Total params: 32,772 (128.02 KB)

Trainable params: 32,772 (128.02 KB)

Non-trainable params: 0 (0.00 B)

Jena classification task

In order to solve the Jena classification problem, a 4-neurone output layer was employed as part of the softmax activation instead of using the regression output layer. These 4 output neurones represent each class of temperature, namely cold, cool, mild, and warm. A softmax activation function works well in this case since the resulting class predictions sum up to one. The loss function was adjusted from mean absolute error to categorical cross-entropy.

b. Train and evaluate: accuracy, macro-F1, confusion matrix

TRAIN LSTM CLASSIFICATION MODEL

Epoch 1/50

200/200 - 21s - 107ms/step - accuracy: 0.5878 - loss: 0.9596 - val_accuracy: 0.7891 - val_loss: 0.7891

Epoch 2/50

200/200 - 19s - 94ms/step - accuracy: 0.6497 - loss: 0.8079 - val_accuracy: 0.7734 - val_loss: 0.7734

Epoch 3/50

200/200 - 20s - 98ms/step - accuracy: 0.6611 - loss: 0.7720 - val_accuracy: 0.7188 - val_loss: 0.7188

Epoch 4/50

200/200 - 19s - 94ms/step - accuracy: 0.6727 - loss: 0.7427 - val_accuracy: 0.8594 - val_loss: 0.8594

Epoch 5/50

200/200 - 19s - 96ms/step - accuracy: 0.6702 - loss: 0.7428 - val_accuracy: 0.6406 - val_loss: 0.6406

Epoch 6/50

```
200/200 - 19s - 96ms/step - accuracy: 0.6846 - loss: 0.7193 - val_accuracy: 0.8125 - val_loss: 0.7193
Epoch 7/50
200/200 - 19s - 94ms/step - accuracy: 0.6891 - loss: 0.7006 - val_accuracy: 0.4844 - val_loss: 0.7006
```

CREATE TEST CLASSIFICATION ARRAYS

TEST PREDICTION

```
267/267 - 4s - 15ms/step
```

ACCURACY, MACRO-F1, CONFUSION MATRIX

	Model	Dataset	Test_Accuracy	Macro_F1
1	LSTM	Jena 4-Class	0.6579	0.4475

	Predicted				
Actual	Cold	Cool	Mild	Warm	
Cold	126	641	0	0	
Cool	265	3118	287	0	
Mild	8	648	2314	22	
Warm	0	17	1035	63	

The LSTM classification model

The accuracy attained by the LSTM classifier on testing is 0.6579 while the macro F1-score is 0.4475 . The confusion matrix gives the distribution of classification results for the four temperature classes. Misclassification should happen between adjacent temperature classes, e.g., cool and mild or mild and warm, since their temperature ranges overlap.

c. Discussion: Regression or classification advantage?

As compared to the derived classification problem, the LSTM network will be more advantageous for the Jena regression prediction. The reason is that the LSTM network can take advantage of the continuous temporal dependencies, day-night cycles, and even multiple-day weather patterns in the training data due to its ability to predict the 24-hour-ahead temperature value exactly.

In the case of classification, the continuous value of the target is discretised into four categories. In doing so, part of the information is lost since the temperatures that are close to each other but on the border of two adjacent classes are put into two different categories. For instance, 9.9°C corresponds to the category ‘Cool’, while 10.0°C belongs to the category ‘Mild’.

In theory, the LSTM networks are designed to model sequential dependence thanks to their gating mechanism and hence have an advantage in tasks requiring long-term temporal information extraction. Nevertheless, upon discretisation, some part of the sequential information might get lost, and thus, the LSTM network would be advantageous only for regression, limiting its effectiveness in classification tasks where temporal patterns are crucial.

III.D LSTM for MNIST Sequential Image

a. Reshape MNIST images to (28, 28)

```
[1] 50000      28      28
```

```
[1] 10000      28      28
```

```
[1] 10000      28      28
```

Reshape MNIST images to (28, 28)

MNIST images were reshaped such that they have 28 time steps, whereby each of the 28 time steps represents a row within the image that has 28 pixel features. This means that each image has an input shape of (28, 28).

b. Design and train LSTM model

```
Model: "sequential_9"
```

Layer (type)	Output Shape	Param #
lstm_4 (LSTM)	(None, 64)	23,808
dense_19 (Dense)	(None, 128)	8,320
dense_20 (Dense)	(None, 10)	1,290

Total params: 33,418 (130.54 KB)
Trainable params: 33,418 (130.54 KB)
Non-trainable params: 0 (0.00 B)

COMPILE MNIST LSTM

TRAIN MNIST LSTM

Epoch 1/20

391/391 - 7s - 17ms/step - accuracy: 0.6810 - loss: 0.9411 - val_accuracy: 0.9049 - val_loss

Epoch 2/20

391/391 - 5s - 12ms/step - accuracy: 0.8900 - loss: 0.3439 - val_accuracy: 0.9456 - val_loss

Epoch 3/20

391/391 - 5s - 12ms/step - accuracy: 0.9284 - loss: 0.2298 - val_accuracy: 0.9594 - val_loss

Epoch 4/20

391/391 - 5s - 12ms/step - accuracy: 0.9440 - loss: 0.1799 - val_accuracy: 0.9619 - val_loss

Epoch 5/20

391/391 - 5s - 12ms/step - accuracy: 0.9543 - loss: 0.1468 - val_accuracy: 0.9728 - val_loss

Epoch 6/20

391/391 - 5s - 12ms/step - accuracy: 0.9609 - loss: 0.1279 - val_accuracy: 0.9774 - val_loss

Epoch 7/20

391/391 - 5s - 13ms/step - accuracy: 0.9640 - loss: 0.1160 - val_accuracy: 0.9787 - val_loss

Epoch 8/20

391/391 - 5s - 13ms/step - accuracy: 0.9684 - loss: 0.1021 - val_accuracy: 0.9797 - val_loss

Epoch 9/20

391/391 - 5s - 12ms/step - accuracy: 0.9717 - loss: 0.0924 - val_accuracy: 0.9811 - val_loss

Epoch 10/20

391/391 - 5s - 13ms/step - accuracy: 0.9719 - loss: 0.0866 - val_accuracy: 0.9836 - val_loss

Epoch 11/20

391/391 - 5s - 12ms/step - accuracy: 0.9750 - loss: 0.0792 - val_accuracy: 0.9841 - val_loss

Epoch 12/20

391/391 - 5s - 12ms/step - accuracy: 0.9769 - loss: 0.0738 - val_accuracy: 0.9836 - val_loss

Epoch 13/20

391/391 - 5s - 12ms/step - accuracy: 0.9779 - loss: 0.0716 - val_accuracy: 0.9838 - val_loss

Epoch 14/20

391/391 - 5s - 13ms/step - accuracy: 0.9800 - loss: 0.0627 - val_accuracy: 0.9864 - val_loss

Epoch 15/20

391/391 - 5s - 12ms/step - accuracy: 0.9812 - loss: 0.0610 - val_accuracy: 0.9859 - val_loss

Epoch 16/20

391/391 - 5s - 12ms/step - accuracy: 0.9818 - loss: 0.0571 - val_accuracy: 0.9865 - val_loss

```
Epoch 17/20
391/391 - 5s - 12ms/step - accuracy: 0.9822 - loss: 0.0564 - val_accuracy: 0.9873 - val_loss
Epoch 18/20
391/391 - 5s - 12ms/step - accuracy: 0.9832 - loss: 0.0532 - val_accuracy: 0.9865 - val_loss
Epoch 19/20
391/391 - 5s - 12ms/step - accuracy: 0.9832 - loss: 0.0517 - val_accuracy: 0.9879 - val_loss
Epoch 20/20
391/391 - 5s - 12ms/step - accuracy: 0.9848 - loss: 0.0484 - val_accuracy: 0.9880 - val_loss
```

c. Test accuracy and comparison with CNN

```
313/313 - 1s - 3ms/step - accuracy: 0.9879 - loss: 0.0360
```

```
$accuracy
[1] 0.9879
```

PARAMETER COUNT

```
[1] 33418
```

COMPARE LSTM WITH CNN

	Model	Test_Accuracy	Parameters	Train_Time_Per_Epoch_Seconds
1	2D-CNN	0.9894	1625866	37.14
2	LSTM Sequential Image	0.9879	33418	4.99

The MNIST LSTM model

The MNIST LSTM model was able to get an accuracy of 0.98 on the test set. For the MNIST LSTM model, each image was considered as a sequence of 28 rows, where each row has 28 pixel features.

Since the MNIST handwritten digit classification task involves pattern recognition based on spatial attributes of the images, we expect that the LSTM will be inferior to the 2D-CNN since it treats each input sample as sequences of data points rather than a multidimensional tensor. In particular, the 2D-CNN architecture allows extracting information about local spatial features such as edges and curves because the filters scan the image in two dimensions.

It can be concluded that the accuracy gap between the CNN and the LSTM is due to different types of inductive biases. While the LSTM might have sufficient capacity to classify the digits, the inductive bias used for learning tasks from sequential data is suboptimal in the case of image processing.

Part IV.1 Design Brief — Custom Hybrid Architecture

The chosen custom hybrid architecture is the CNN-LSTM model for solving the temperature regression issue in Jena. The custom hybrid architecture includes a 1D convolutional neural network along with the LSTM model. The reasons motivating the choice of this architecture lie in the weaknesses of the respective standard models.

Specifically, the DNN model uses a fully connected structure which considers all observations in the input dataset as features for regression in the 120-hour period. In other words, such architecture breaks the order of climate observations, which makes it unable to identify local time-dependent relationships. The 1D-CNN (one-dimensional convolutional neural network) solves this issue by detecting short-term time-dependent patterns such as variations of temperature, humidity, and air pressure during one hour. At the same time, this architecture fails to identify long-term dependencies in climate data. The ability to do so belongs to the LSTM networks. On the other hand, they may become computationally expensive and be less efficient in extracting local features due to their nature, which can hinder their performance in real-time applications where quick processing is essential.

The hybrid CNN-LSTM architecture is designed to overcome these issues. First, Conv1D layers enable the extraction of local temporal features out of the multivariate climate sequences. Then, the features obtained are passed to the LSTM layer for building the dependencies between the elements in this sequence. As a result, the custom hybrid architecture incorporates two inductive biases: the extraction of local features through convolution and longer sequential relationships identification through LSTM.

The chosen problem domain of weather forecasting is a good match for the suggested hybrid architecture because weather data include different patterns of temporal behaviour, both short-term ones taking place over a few hours and longer daily patterns covering a 120-hour period. Since the target variable of regression is the future 24-hour temperature value, a model that considers both local changes in the climate data and its long-term dependencies would work properly.

IV.1 Design Brief (Using Suggested Hybrid Pattern: CNN–LSTM)

A hybrid architecture used in this work is CNN-LSTM, which consists of one-dimensional convolutional layers together with LSTM layers to solve Jena’s temperature regression problem.

The reason for choosing such architecture lies in the shortcomings of pure models like DNN, CNN, or LSTM.

First of all, the DNN-based model disregards the temporal nature of the data, as it uses the flattened representation of the input data. The second model, 1D-CNN, performs better compared to DNN due to its ability to use filters for capturing the local temporal dependencies in the data. However, it cannot model dependencies between observations located far from each other in the sequence of the 120-hour lookback window. Contrarily, LSTM is able to learn long-range dependencies due to its gating function; however, it might be less efficient at extracting short-range local patterns than convolutions, which can lead to challenges in accurately modelling data that requires both short and long-range context.

The combination of CNN and LSTM resolves both shortcomings of models mentioned above, as the CNN-based feature extraction layer extracts local temporal features such as local weather changes (hourly, daily, or week-based). Thereafter, the sequence of features is passed as input data to an LSTM-based layer, which learns long-term dependencies and temporal relationships. Consequently, such a network learns two inductive biases: local features using CNN and long-range memory using LSTM.

The Jena climate dataset is well-suited for testing of the described approach since the weather is characterised by periodicity in the form of weather changes. Moreover, weather changes can also occur locally within a period. Hence, since we aim to make a forecast in the next 24 hours based on 120-hour history, our architecture is suitable for such a task.

IV.2 CNN-LSTM HYBRID MODEL USING FUNCTIONAL API

Model: "CNN_LSTM_Hybrid_Jena_Regression"

Layer (type)	Output Shape	Param #
Input_120_hours (InputLayer)	(None, 120, 14)	0
Conv1D_Local_Features_1 (Conv1D)	(None, 120, 32)	2,272
Conv1D_Local_Features_2 (Conv1D)	(None, 120, 64)	10,304
MaxPooling_Temporal_Reduction (MaxPooling1D)	(None, 60, 64)	0
LSTM_Long_Range_Dependencies (LSTM)	(None, 64)	33,024
Dense_Feature_Integration (Dense)	(None, 64)	4,160

Temperature_Output (Dense) (None, 1) 65

Total params: 49,825 (194.63 KB)
Trainable params: 49,825 (194.63 KB)
Non-trainable params: 0 (0.00 B)

COMPILE HYBRID MODEL

TRAIN HYBRID MODEL

Epoch 1/50
200/200 - 15s - 75ms/step - loss: 0.3384 - mae: 0.3384 - val_loss: 0.4929 - val_mae: 0.4929
Epoch 2/50
200/200 - 10s - 48ms/step - loss: 0.2912 - mae: 0.2912 - val_loss: 0.5133 - val_mae: 0.5133
Epoch 3/50
200/200 - 10s - 48ms/step - loss: 0.2857 - mae: 0.2857 - val_loss: 0.3823 - val_mae: 0.3823
Epoch 4/50
200/200 - 10s - 50ms/step - loss: 0.2742 - mae: 0.2742 - val_loss: 0.3659 - val_mae: 0.3659
Epoch 5/50
200/200 - 10s - 48ms/step - loss: 0.2730 - mae: 0.2730 - val_loss: 0.3482 - val_mae: 0.3482
Epoch 6/50
200/200 - 10s - 48ms/step - loss: 0.2688 - mae: 0.2688 - val_loss: 0.3387 - val_mae: 0.3387
Epoch 7/50
200/200 - 10s - 49ms/step - loss: 0.2644 - mae: 0.2644 - val_loss: 0.4024 - val_mae: 0.4024
Epoch 8/50
200/200 - 10s - 51ms/step - loss: 0.2621 - mae: 0.2621 - val_loss: 0.3658 - val_mae: 0.3658
Epoch 9/50
200/200 - 10s - 48ms/step - loss: 0.2563 - mae: 0.2563 - val_loss: 0.3551 - val_mae: 0.3551
Epoch 10/50
200/200 - 10s - 48ms/step - loss: 0.2557 - mae: 0.2557 - val_loss: 0.3490 - val_mae: 0.3490
Epoch 11/50
200/200 - 10s - 48ms/step - loss: 0.2526 - mae: 0.2526 - val_loss: 0.3823 - val_mae: 0.3823

Architecture Diagram (CNN-LSTM Hybrid Architecture Diagram)

Input Layer
(120, number of predictors)

Conv1D Layer 1
32 filters, kernel size = 5, ReLU, same padding

Conv1D Layer 2
64 filters, kernel size = 5, ReLU, same padding

MaxPooling1D
pool size = 2

LSTM Layer
64 units, recurrent dropout = 0.20

Dense Layer
64 neurons, ReLU

Output Layer
1 neuron, linear activation

Implementation of the CNN-LSTM Hybrid Architecture

The CNN-LSTM hybrid model was built through the functional API provided by Keras. It fulfils the requirement of designing a hybrid model architecture which goes beyond the traditional sequential model architecture and is created explicitly using `layer_input` and `layer_connections`.

The model receives as input a tensor with dimensions (120, number of predictors), corresponding to the previous 120 hours of Jena climate variables time series. The first two Conv1D layers extract local features in the temporal domain of the input sequence. The selected kernel size is 5 because the filters will learn features corresponding to short five-hour weather patterns. After applying the Conv1D layers, the temporal dimension is downsampled, and only relevant local information remains through the MaxPooling1D layer.


The generated sequence of features is fed into an LSTM layer with 64 units, which can learn temporal dependencies after extracting the short-term patterns. Finally, the sequence of features is aggregated by a dense layer and outputted by a single linear neurone which performs the prediction of the 24-hour-ahead temperature.

The model was trained through the same chronologically defined data splits and early stopping criteria applied to the other Jena models using the MAE loss function.

IV.3 HYPERPARAMETER: NUMBER OF FILTERS

	Filters	Validation_MAE
1	16	0.3724
2	32	0.3781
3	64	0.3509

Plot Hyperparameter Sensitivity



Hyperparameter Sensitivity Analysis

The hyperparameter chosen for sensitivity analysis is the number of convolutional filters. Three models of CNN-LSTM hybrid networks were trained by varying the number of filters to 16, 32, and 64 in the first layer.

As seen from the graph above, the effect of increasing the number of filters on the validation MAE is presented. With the fewest number of filters, the model might suffer from underfitting, since it will be lacking in complexity to capture useful temporal patterns. With more filters, the model has better extraction abilities. At the same time, too much complexity might increase the risk of overfitting or high computational costs.

The plot illustrates that there is an optimal range of filter values when validation MAE is minimised. This demonstrates that the hybrid model is sensitive to the choice of convolutional filters, and careful tuning is required to achieve the best performance.

HYBRID MODEL TEST EVALUATION

268/268 - 2s - 7ms/step

	Model	Test_MAE	Test_RMSE	Test_R2
1	CNN-LSTM Hybrid	0.2772	0.3559	0.8611

Evaluation and Reflection

The results for the hybrid CNN-LSTM model were that it had a test MAE of 0.2772, RMSE of 0.3559, and R^2 of 0.8611 when tested on the Jena data set. This shows the accuracy of the model in predicting the temperature after 24 hours.

Since the hybrid CNN-LSTM model utilises the benefits of both the convolutional neural network architecture and the recurrent neural network architecture, we would expect it to be better than the individual networks, which include DNN, 1D-CNN, and LSTM.

If the hybrid model outperforms the individual models, then this validates that the fusion of multiple inductive biases enhances its performance. If the hybrid model underperforms the individual models, one plausible reason could be that the added complexity of the hybrid model causes issues such as overfitting and optimisation problems. Another reason for underperformance could be that the individual models have captured most of the signal in the dataset.

Consequently, the performance of the hybrid model depends on its complexity and the nature of the data.

Part V.1 REGRESSION SUMMARY TABLE

	Model	Test_MAE	Test_RMSE	Test_R2	Parameters
1	DNN	0.3016	0.3848	0.8377	223489
2	1D-CNN	0.2915	0.3712	0.8490	16801
3	LSTM	0.2699	0.3468	0.8682	32673
4	CNN-LSTM Hybrid	0.2772	0.3559	0.8611	49825
	Train_Time_Per_Epoch_Seconds				
1	NA				

2	10.73
3	19.61
4	10.21

V.1 Classification Models Summary Table

	Model	Dataset	Test_Accuracy	Macro_F1	Parameters
1	DNN	MNIST	NA	NA	NA
2	2D-CNN	MNIST	0.9894	0.9893	1625866
3	LSTM (MNIST)	MNIST	0.9879	NA	33418
4	LSTM (Jena)	Jena 4-Class	0.6579	0.4475	32772
	Train_Time_Per_Epoch_Seconds				
1	NA				
2	37.14				
3	4.99				
4	NA				


Summary of Model Performance

In the regression table, the performance of the Deep Neural Network (DNN), one-dimensional Convolutional Neural Network (1D-CNN), Long Short-Term Memory (LSTM), and CNN-LSTM hybrid models is assessed against the Jena temperature prediction problem. In addition to measures like Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and coefficient of determination (R^2), model size (parameter count) and computational complexity (epoch training time) are also considered.


The classification summary table evaluates the models' performance based on accuracy and macro-F1 score between the MNIST and Jena datasets. It is anticipated that the CNN will yield the highest accuracy for the MNIST dataset since it can learn spatial characteristics. In the same vein, the LSTM for MNIST is anticipated to be less accurate compared to the CNN since it views the images as sequences instead of spatial entities. On the other hand, when it comes to the Jena dataset, the LSTM can capture temporal dependencies; however, it will not yield as much improvement as in the case of the regression since there will be some loss in information, particularly because the LSTM's focus on sequences may overlook certain spatial features present in the data.

In conclusion, the above summary tables present a detailed analysis of the model performance.

V.2 Visualizations Plot: Regression Predictions vs Actual (Two Weeks)



Bar Chart: Classification Accuracy



Visualization of Model Performance

The regression plot shows the comparison between the estimated temperatures using different models with the real temperature readings within a span of two weeks. Hybrid CNN-LSTM model and LSTM model would be able to mimic the true temperature pattern better than the other models, particularly where there is drastic temperature variation.

The classification accuracy bar graph shows the comparative results for the accuracy in MNIST and Jena datasets. Two-dimensional CNN gives the best classification accuracy in the MNIST dataset because of its spatial feature, while LSTM is competitive as per the structure of the dataset.

V.3 Statistical Comparison (Diebold–Mariano Test)

Diebold–Mariano Test

```
data: errors_lstmerrors_hybrid
```

DM = -5.4131, Forecast horizon = 1, Loss function power = 1, p-value = 6.363e-08
alternative hypothesis: two.sided

Statistical Comparison of Forecast Accuracy

A comparison of the forecast accuracies of the two best regression models is undertaken through a Diebold-Mariano (DM) test. This test helps in determining whether the difference in the error prediction between the two models is statistically significant.

In the DM test, the null hypothesis assumes that the two regression models exhibit equal predictive accuracy, while the alternative hypothesis assumes otherwise.

When the p-value is 0.05, the null hypothesis is rejected, implying that the difference in MAE between the two models is statistically significant. Conversely, when the p-value is 0.05, the difference between the two models is statistically insignificant; thus, both models can be used for forecasting.

This statistical tool is important in model comparison and selection because it allows us to make an informed decision by supporting our claims using statistics.

Critical Reflection

1. Best Architecture for Regression and Theoretical Justification

The CNN-LSTM model (or the LSTM model, depending on your experimental results) worked better for Jena regression. This is not surprising in light of theoretical considerations, since predicting the temperature has both short- and long-term dependencies. The LSTM model includes inductive bias to work with sequential information by means of its memory cells and gates, thus being able to identify long-term dependencies within the 120-hour span of time. The CNN-LSTM model improves upon that result by first identifying local dependencies over time using convolutional filters and then modelling long-range dependencies using LSTMs.

2. Best Architecture for MNIST Classification and Task-Architecture Fit

The 2D-CNN model had superior performance when applied to the MNIST classification problem. This is due to the high level of correlation between the architecture and the problem since the images in the MNIST dataset have spatial characteristics. The CNN leverages this characteristic using the edge detectors, curve detectors, and shape detectors within the convolution operation. It further utilises parameter sharing and spatial pooling. On the other hand, the LSTM network considers the image as a series of rows and thus cannot leverage all the two-dimensional spatial correlations in the image.

3. Universal Approximation Theorem vs Practical Performance

The Universal Approximation Theorem claims that with a sufficiently large network, it is possible to approximate any function. Still, the outcomes illustrate that the DNN fails to reach the same level of efficiency as either CNN or LSTM models. It is due to the lack of proper inductive biases for structured data like image and sequence data in the DNN architecture. Although it is capable of capturing those patterns, it takes much more effort, data, and training time compared to CNNs and LSTMs. The latter architectures, such as Convolutional Neural Networks (CNNs) for image data and Long Short-Term Memory networks (LSTMs) for sequence data, are built to capitalise on structural properties from the very beginning, which makes them both more efficient and effective in solving tasks compared to traditional Deep Neural Networks (DNNs).